

Ishaan Jhaveri

Instructor: Amy Williams

BTRY 4840/6840, CS 4775 Fall 2016: Computational Genetics and Genomics

14 December, 2016

## What can alleles on certain Single Nucleotide Polymorphisms (SNPs) with anomalous p-values calculated using Gene Dropping (GD) techniques tell us about genetic variation in a natural population?

### 1 Introduction

Genotypes of members of any population over time are influenced by genetic forces. Forces such as inbreeding lead to more homogeneous genotypes, whereas forces such as mutation lead to heterogeneous genotypes. Other genetic forces like natural selection are also at play. An interesting problem would be to calculate and quantify just how much these forces contribute to genetic variation in a population. To this end, Gene Dropping (GD) is a popular technique used by population geneticists to solve this problem. The idea is use computer simulations to create a null hypothesis for what the genotypes of members a population would look like if all individuals bred completely randomly, and all alleles passed from parent to offspring purely following Mendel's Law of Independent Assortment. The null hypothesis for each individual's genotype can then be compared to each individual's actual genotype. Differences in the simulated genotypes and actual genotypes can give insights into the influence of various genetic forces. In order to do this one needs a pedigree of individuals with genotypic data for each individual, as well as information about how each individual is related to the other individuals. In this project, we create a program that can perform GD, and test it on data from an endemic population of Florida Scrub-Jay's consisting of a little over twelve thousand individuals. The genotypic data for these individuals is at the Single Nucleotide Polymorphism level. In the test dataset we have 1262 SNPs for each individual.

#### 1.1 Related Work

Macluer et al (1986) used GD to analyze the extent of genetic variability in a colony of small South American marsupials, *Monodelphis domestica*. Since then GD programs like Mendel, Merlin, GENLIB, IBDsim, LDSO and others have been used by many researchers for many ends. To name a few, Gao et al used GD techniques to Estimate of the Average Number of Recessive Lethal Mutations Carried by Humans (2015). Kennedy et al used GD to measure the rate of Inbreeding Depression among the Chatham Island Black Robin (2013).

In this project in particular we have followed Macluer et al's method very closely. We have included Figures 1 and 2 below from Macluer's paper. As per Macluer's method, GD begins with the assignment of two unique hypothetical alleles to each of the founding members of a colony. Figure 1 illustrates the process for a simple pedigree consisting only of animal A and his ancestors. A random number generator is used to determine which one of each founders genes is transmitted to each offspring. At the end of a single gene drop, every animal in the pedigree has a genotype (Fig. 2).




Fig. 1. Pedigree with two unique hypothetical alleles assigned to each founder.

Figure 1: Figure 1 from Macluer et al (1986).

Macluer's method performed thousands of drops at a time. The method in this project aims to be used on the order of several hundred thousand to a million drops, to generate richer data. The program developed through the course of this project is also highly customizable, as an improvement on Macluer's original implementation.

## 2 Work with Nancy Chen and the Clark Lab

The purpose for this project came about in the Summer of 2015. Nancy Chen, at the Clark Lab of Cornell University was conducting research that would be used to write a paper that tackles the question, "what is the importance of recent natural selection in governing patterns of genetic variation in a natural population?".

This project was seeking a customizable, fast implementation of Gene Dropping. The author and a fellow undergraduate studying Computer Science wrote a program in Python to simulate the GD Chen would need to answer this question. The developed program takes in a pedigree file, which encodes the necessary information as explained above. It also takes in a file mapping breeding pairs to the year in which they bred and a file mapping SNPs to the chromosomes they occur on (format explained in Appendix B). The program then "drops" genes down the pedigree, using Macluer's random approach. For each individual it randomly selects an allele at each SNP from each parent. Special cases are delineated further in the section on Methods. This program was written in Python.




Fig. 2. Pedigree after gene dropping, with genotypes assigned to descendants by Mendelian segregation of founder alleles.

Figure 2: Figure 2 from Macluer et al (1986).

## 2.1 Issues and Shortcomings of this work

The program was written in Python for its usability and wealth of libraries for file manipulation and data handling. However Python's speed proved to be prohibitively expensive. The Python program could only perform at about 30 seconds a drop which amounts to almost 4 days to run 10,000 gene drops and more than a month to run 100,000 gene drops. Since this is the lower end of the magnitude of drops this program would be used for, this speed rendered the program useless for the purpose of this research.

Note: All drop time measurements were taken on a computer with 2.6 GHz processing power.

## 3 Extension: Speedup

Thus the primary purpose of this project is a speedup on the method developed by the author and others in the Clark Lab in 2015.

### 3.1 Method for achieving speedup

The Python implementation of the GD program is very robust in terms of preprocessing the data. Thus for this project we will use some of the preprocessing logic already written into the Python implementation. However preprocessing is only a one-time requirement, so it takes a constant time, after which, with clean enough data, a C program can perform the actual drops down the pedigree much faster than the Python dropping implementation.

Therefore the scope of this project was making use of the existing Python implementation to preprocess the data and writing a C implementation to perform the actual dropping logic. A necessary last step is displaying the results of the C implementation's drops, which is also in the scope of this project.

## 4 Method

This section will explain the entire GD implementation from preprocessing to displaying the data.

### 4.1 Preprocessing

The first step in Preprocessing is taking isolated individuals out of the pedigree. Isolated individuals are defined as individuals that appear in the pedigree that have no parents and no offsprings, either due to lack of data, or for any other reason.

The next step is dealing with founders. Founders are defined as the breeding pairs of individuals in the first breeding year under scrutiny. The first breeding year can either be input directly (eg: 2001), or calculated from the data. If the user selects the calculation from data option, the program chooses the first year as the year in which there is genotype data for at least a certain fraction (called the dense data threshold - also user specified) of breeding pairs in that year. The breeding pairs in the first year that meets this criteria are chosen as "founders". Founders are important because the founding generation's allele frequencies are treated as the baseline allele frequencies in data analysis. Founders' genotypes are the ones dropped throughout the pedigree, so these are the genotypes we care most about.

Therefore the program has settings built in to deal with missing founder genotype data. The user can select one of four options - impute, simulate, leave unknown, user specified. Impute means that if a founder's genotype data is missing, it is imputed from the genotypes of its known offsprings. Simulate means the founder is assigned likely genotypes based on the allele frequencies of other founders' genotypes. Leave unknown, is as the name suggests an option to leave unknown founders' genotypes unknown and drop an 'X' down the pedigree to denote the genes of this founder. User specified allows the user to provide one line of a pedigree file that has genotype data for all missing founders' genotype data.

The last step in preprocessing is dealing with immigrants. Immigrants are defined as individuals with offsprings that either have no parents, or have no parents that bred in the founding generation. These individuals are also interesting because they add genes to the population as the gene dropping continues, but they are not founders, so they might confuse the data analysis. In order to maintain the sanctity of the data analysis, the user also has the following options on how to deal with immigrants. The user can select one of five options - impute, simulate, leave unknown unknown, leave unknown all and user specified. Impute means that if an immigrants's genotype data is missing, it is imputed from the genotypes of its known offsprings. Simulate means the immigrant is assigned likely genotypes based on the allele frequencies of founders' genotypes. Leave unknown unknown, is an option to leave unknown immigrants' genotypes unknown and drop a 'U' down the pedigree to denote the genes of this immigrant. Leave unknown all, is like leave unknown unknown, but all the immigrants' genes (even known ones) are dropped as 'U' down the pedigree. User specified allows the user to provide one line of a pedigree file that has genotype data for all missing immigrants' genotype data.

## 4.2 Dropping and Data Collection

The drop step follows a simple inner-outer loop structure. The inner loop assigns alleles to each SNP of a particular individual based on a random coin flip from their mother's alleles at that SNP and the same for their father's. The outer loop runs through the entire pedigree, so that each individual gets a simulated genotype. It makes sure to simulate genotypes in an order such that when simulating an individuals' genotype, its parents always already have a simulated genotype.

As it assigns each individual a genotype, it also updates a data structure keeping track of the various counts of each allele at each SNP across all the breeders for each breeding year.

## 4.3 Post processing

The post processing consists of printing out the above data structure in the following format.

1. YEAR is the breeding year for which this line gives data. The data will be the information collected from all the breeders in that year.
2. SNP is the SNP id.
3. NUM\_ALLELES is the number of observed alleles at this SNP. This will always be 1 or 2 unless the 'U' or 'X' options are acitvated.
4. ALLELE:COUNT... is each allele and its corresponding count.

5. `TIMES_THESE_COUNTS_OCCURED` is the number of times this configuration of counts occurred with this YEAR and SNP.

The user can also input a list of YEAR, SNP pairs or select 'all' if the user does this, in a separate file the p\_beginning, and p\_observed values for that YEAR, SNP pair are printed in a separate file. With these aspects of the data, the user can calculate p-values as in Figure 3 below. The user also has an option telling the program to generate graphs, in which case the graphs look like the ones shown in the section on Results.

## 5 Improvement on Macluer's work: Customizability

Clearly this program is highly customizable. The various aspects the user can control are as follows:

1. The pedigree to drop on.
2. The number of drops to perform.
3. Whether to calculate the founding year or specify it.
4. Whether to impute, simulate, leave unknown or specify missing founder data.
5. The same for immigrants.
6. Which, if any YEAR, SNP pairs to print more detailed data about.
7. Whether to output graphs or not.

## 6 Results

### 6.1 Results of attempted Speedup

The final speed (on a 2.6 GHz processor) is about 50 seconds for preprocessing, then about 0.2 seconds per drop and then about  $(0.05 \times \text{number of drops})$  seconds to output the data correctly.

### 6.2 Example data generated

Figures 4 and 5 are snapshots of the sample output generated by running 1,000 gene drops on the entire pedigree. The entire file is also included with the submission as `data/one_thousand.out`.

The headings are explained in the subsection on Postprocessing in the section on Methods.

Figure 6 is an example of the p-value data outputted by the program.

1. YEAR is the year in question
2. SNP\_NUM is the index of the SNP in question
3. SNP\_ID is the id of the SNP in question
4. NUM\_ALLELES\_AT\_SNP is the number of alleles seen at this SNP in any year.

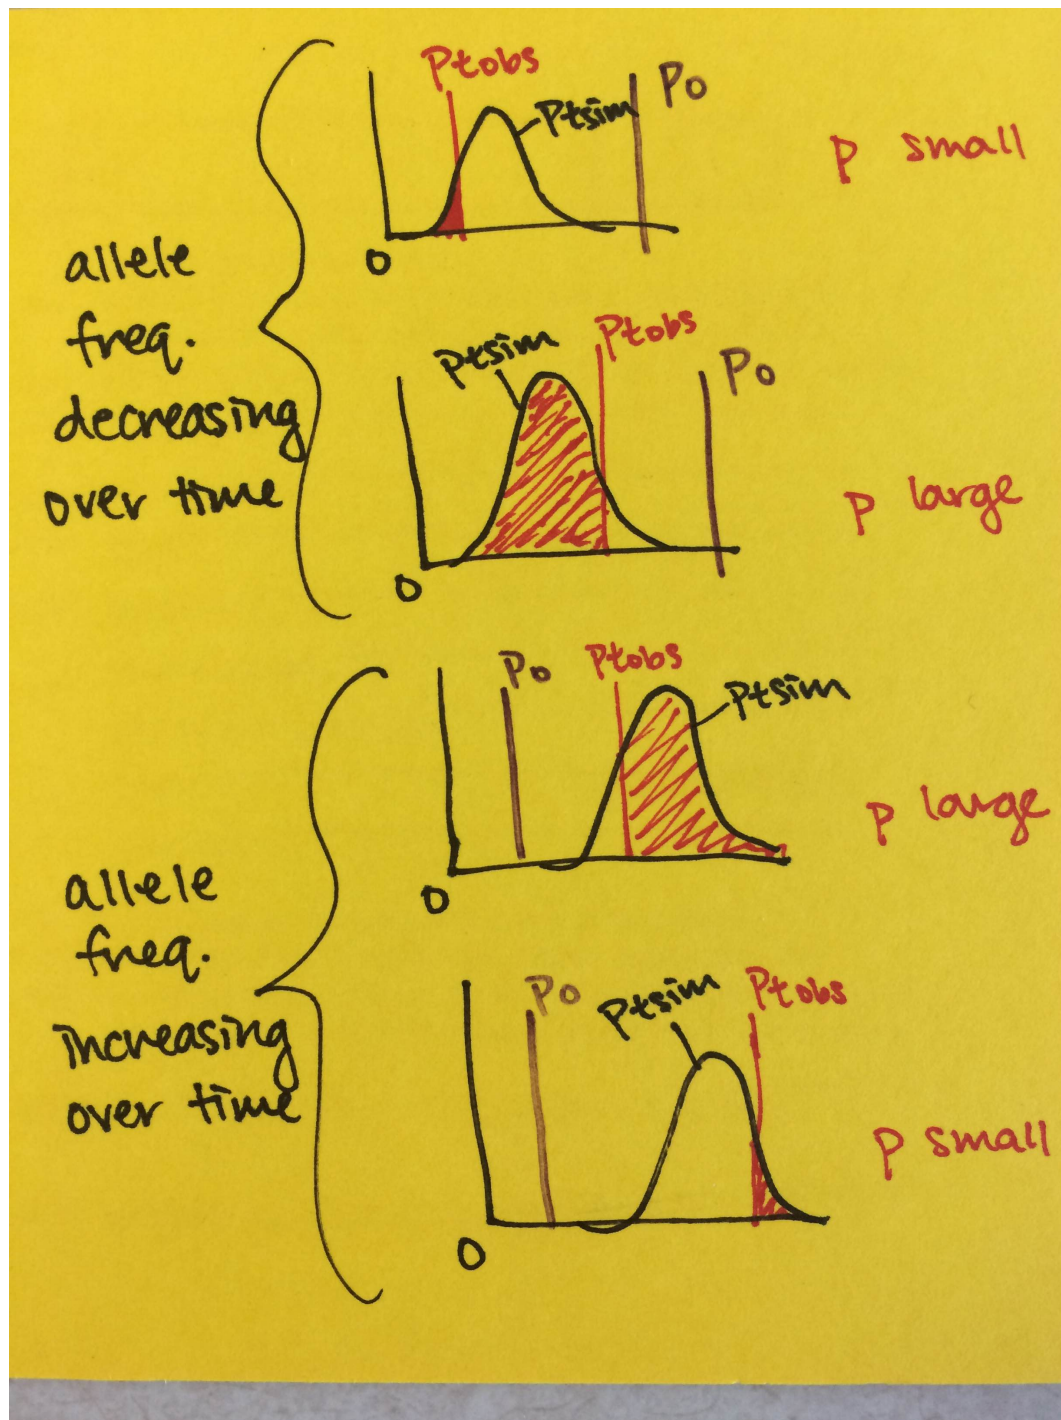


Figure 3: P-value calculation, courtesy of Nancy Chen.

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25
	YEAR	SNP	NUM_ALLELES	ALLELE:COUNT...	TIMES_THESE_COUNTS_OCCURED																			
2	2001	1	2	A:250 C:238	1000																			
3	2001	2	2	C:186 T:302	1000																			
4	2001	3	2	A:216 C:272	1000																			
5	2001	4	2	A:264 G:224	1000																			
6	2001	5	2	A:262 G:226	1000																			
7	2001	6	2	G:303 T:185	1000																			
8	2001	7	2	C:192 T:296	1000																			
9	2001	8	2	C:234 T:254	1000																			
10	2001	9	2	C:181 T:307	1000																			
11	2001	10	2	A:230 G:258	1000																			
12	2001	11	2	C:246 T:242	1000																			
13	2001	12	2	A:284 G:204	1000																			
14	2001	13	2	C:276 T:212	1000																			
15	2001	14	2	A:241 G:247	1000																			
16	2001	15	2	A:150 G:338	1000																			
17	2001	16	2	A:188 G:300	1000																			
18	2001	17	2	C:230 T:258	1000																			
19	2001	18	2	C:229 T:259	1000																			
20	2001	19	2	A:242 G:246	1000																			
21	2001	20	2	C:283 T:205	1000																			
22	2001	21	2	A:255 C:233	1000																			
23	2001	22	2	A:279 G:209	1000																			
24	2001	23	2	C:250 T:238	1000																			
25	2001	24	2	A:202 G:286	1000																			

Figure 4: Example output

5. ALLELES:P\_FIRST\_GENERATION... are these alleles with their frequencies in the founding generation.
6. ALLELES:P\_OBSERVED... are these alleles with their frequencies in the observed generation.

Figure 7 is an example graph generated by the program.

## 7 Scientific question and other applications of this work

Clearly using this data technique one can find out which SNPs show anomalous behaviour, and make various inferences. Since we also have a map from SNPs to chromosomes we can make phenotypic inferences about anomalous genotypic behaviour at particular SNPs as shown by this data.

One can also correlate the phenotypes that correspond to changes in these SNPs and understand what traits might be selected for or evolutionarily beneficial.

This program will now be used by Chen et al to answer the original scientific question they set out to answer, "what is the importance of recent natural selection in governing patterns of genetic variation in a natural population?"

## 8 Extensions and Areas to improve the code

The section of the code to generate graphs and read user specified genotypes for founders or immigrants is a bit buggy as of now, and needs some attention. The graphs' headings and labeling



1249	2001	1248	2	A:214	G:274	1000
1250	2001	1249	2	C:284	T:204	1000
1251	2001	1250	2	A:355	G:133	1000
1252	2001	1251	2	C:137	T:351	1000
1253	2001	1252	2	A:319	G:169	1000
1254	2001	1253	2	A:215	G:273	1000
1255	2001	1254	2	G:192	T:296	1000
1256	2001	1255	2	A:230	C:258	1000
1257	2001	1256	2	A:289	G:199	1000
1258	2001	1257	2	G:261	T:227	1000
1259	2001	1258	2	C:186	T:302	1000
1260	2001	1259	2	C:307	T:181	1000
1261	2001	1260	2	A:225	G:263	1000
1262	2001	1261	2	C:277	T:211	1000
1263	2001	1262	2	A:218	G:270	1000
1264	2002	1	3	A:233	C:212	U:105 17
1265	2002	1	3	A:229	C:217	U:104 16
1266	2002	1	3	A:230	C:218	U:102 10
1267	2002	1	3	A:239	C:205	U:106 1
1268	2002	1	3	A:228	C:218	U:104 8
1269	2002	1	3	A:229	C:213	U:108 7
1270	2002	1	3	A:234	C:212	U:104 16
1271	2002	1	3	A:234	C:213	U:103 16
1272	2002	1	3	A:229	C:212	U:109 2
1273	2002	1	3	A:238	C:208	U:104 6
1274	2002	1	3	A:230	C:214	U:106 10
1275	2002	1	3	A:234	C:210	U:106 7
1276	2002	1	3	A:225	C:220	U:105 1
1277	2002	1	3	A:232	C:218	U:100 5

Figure 5: Example output

YEAR	SNP_NUM	SNP_ID	NUM_ALLELES_AT_SNP	ALLELES:P_FIRST_GENERATION...	ALLELES:P_OBSERVED
2013	1	s9945p2883	2	A:0.525 C:0.475 A:0.56875 C:0.43125	
2015	1	s9945p2883	2	A:0.525 C:0.475 A:0.59375 C:0.40625	

Figure 6: Example line of a p data output file

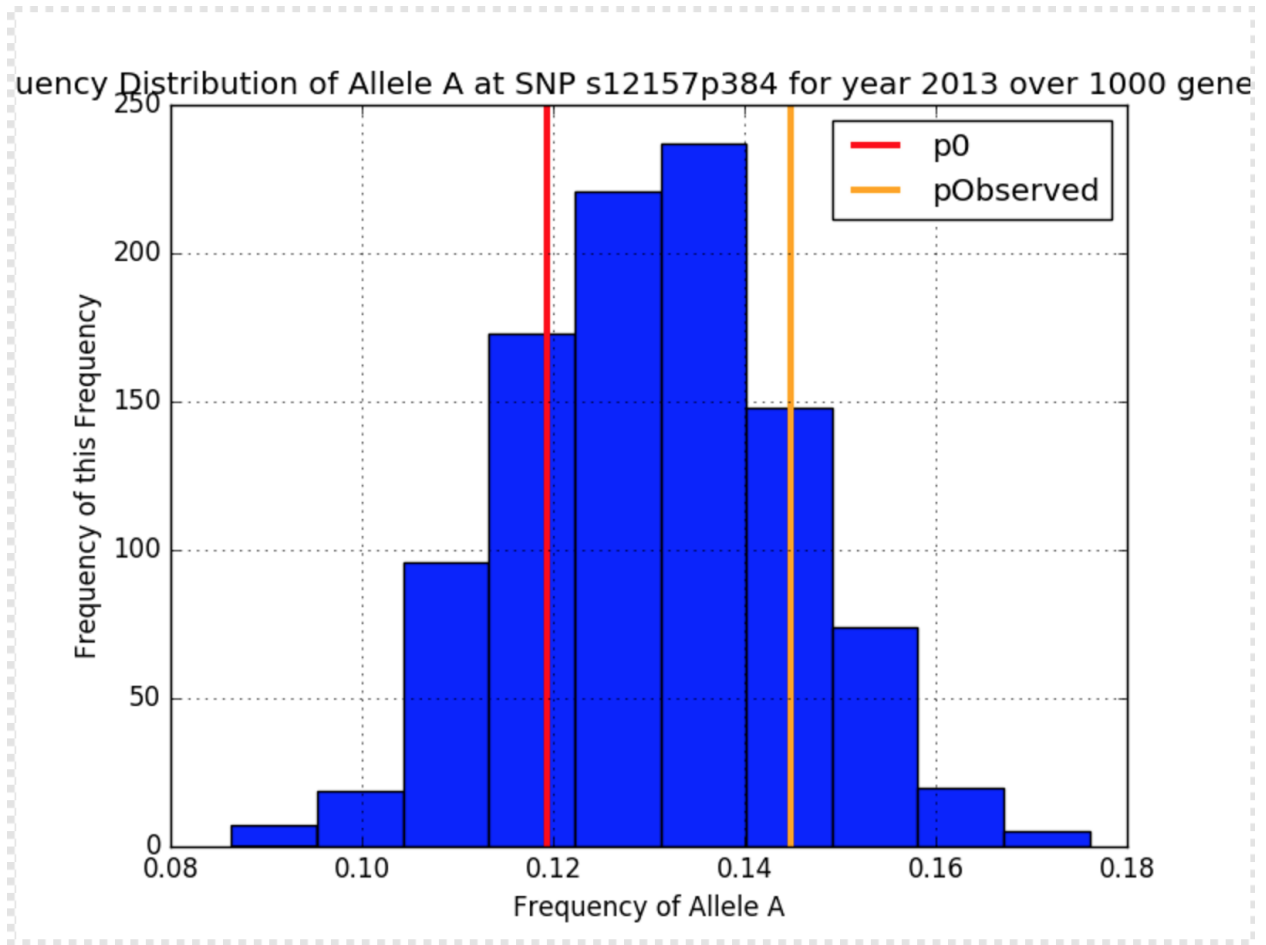


Figure 7: Example graph generated by program from data output and p data output

can be improved too.

As of now the code treats autosomes and allosomes the same, but it can be extended to treat allosomes differently, since genes are not dropped similarly randomly at allosomes. The Python implementation dealt with this, so an extension to the C implementation would be adding functionality to deal with this.

Currently the post-processing, data output phase is not optimized at all. The data is output to a file, and then a Python script scans and cleans that file, collapsing identical lines, and incrementing their count. This step should be done programmatically for enhanced speed.

Another extension is outputting a metadata file, which has information about which and how many individuals were isolated, which and how many founders didn't have genotype data, which and how many immigrants there were etc.

Since we have SNP to chromosome maps, two immediate extensions to this work are local phasing and haplotype inference. Using the observed data and seeing SNPs that behave similarly, one can look at the chromosomes they are on, and make inferences about recombination events.

## 9 References

1. *Pedigree Analysis by Computer Simulation*. MacCluer, VandeBerg, Read and Ryder. 1986. Research Department, San Diego Zoo, San Diego.
2. *An Estimate of the Average Number of Recessive Lethal Mutations Carried by Humans*. Gao, Waggoner, Stephens, Ober and Przeworski. 2015. Columbia University, New York, New York.
3. *Severe inbreeding depression and no evidence of purging in an extremely inbred wild species the chatham island black robin*. Kennedy, Grueber, Duncan and Jamieson. 2013. Institute for Applied Ecology, University of Canberra, Australia.
4. [https://en.wikipedia.org/wiki/Mendelian\\_inheritance](https://en.wikipedia.org/wiki/Mendelian_inheritance). For information on Mendel's laws.

I would like to thank Nancy Chen and Andrew Clark for introducing me to this project. I would like to thank Ashish Akshat for co-writing the first attempt at the project with me.

## 10 Appendices

### 10.1 Appendix A: Configuring and running the code

The code is included as `geneDropping.zip`. Unzipping this directory and navigate to it in a Terminal session.

1. type `cd src`. Hit enter.
2. Edit the following files:
  - 2.1. `constants.py`:

- 2.1.1. DENSE\_DNA\_THRESHOLD: Edit to desired value. This is the fraction of breeders to have data for a year to be considered the founding year.
- 2.1.2. FOUNING\_YEAR: Can specify user defined founding year or leave as 0.
- 2.1.3. IMPUTE\_IMM: True or False depending on whether you want to impute immigrant alleles or not.
- 2.1.4. IMM: 'sim', 'ua', 'uu' or filename for the simulate, unknown unknown, unknown all and specify options explained above.
- 2.1.5. IMPUTE\_FOUNDEERS: True or False depending on whether you want to impute founder alleles or not.
- 2.1.6. FOUNDER: 'sim', 'x' or filename for the simulate, unknown and specify options explained above.
- 2.1.7. P\_VALUES\_FOR: list of YEAR SNP pairs. SNPs can be specified either as indices or as string SNP ids. Can be blank.
- 2.1.8. P\_GRAPH: True if the program should output a graph, False otherwise. Invalid if P\_VALUES\_FOR list is empty.
- 2.1.9. PEDFILE: The pedigree file.
- 2.1.10. MAPFILE: The file mapping SNPs to Chromosomes.
- 2.1.11. BREEDERFILE: The file mapping years to the pairs of individuals that were breeders in that year.
- 2.2. geneDrop.h:
  - 2.2.1. gene\_length: the number of SNPs.
  - 2.2.2. founding\_year: the founding year, if user specified. Leave 0 otherwise.
  - 2.2.3. N: the number of drops to simulate.
- 3. type `chmod 755 runGeneDropping`. Hit Enter.
- 4. type `./runGeneDropping`. Hit Enter. Note this will print the data onto the terminal, to pipe the output to a file type `./runGeneDropping > filename.ext`.

## 10.2 Appendix B: Input File formats

PED FILE: Figure 8.

1. "1" is the family ID.
2. "1013-63348" is the individual ID.
3. "1053-14500" is the father's ID in the ZW system and the mother's in the XY system.
4. "1423-42714" is the mother's ID in the ZW system and the father's in the XY system.
5. "2" is the sex, in this case Male.
6. "-9" is phenotypic information.
7. "C A ..." is the genotype

1 1013-63348 1053-14500 1423-42714 2 -9 C A T T A A G G G G T G

Figure 8: Example line of a .ped file

1 s6401p16442 0 404959

Figure 9: Example line of a .map file

MAP FILE: Figure 9

- 1. "1" is the chromosome number.
- 2. "s6401p16442" is the SNP id.

The other two lines are ignored for the purposes of this project.

BREEDER FILE: Figure 10. We are interested in "Year", "MBreederUSFWS" and "FBreederUSFWS"

TerrYr	Terr	Year	Tract	PairNo	MBreederUSFWS	MBreederBands	FBreederUSFWS	FBreederBands
ANGB08	ANGB	2008	Demo	1	1573-97942	ACH-	1603-16663	-WCG

Figure 10: Example line of the breeder mapping file